

SQL II

R & G - Chapter 5



SELECT DISTINCT



- `SELECT DISTINCT (col list)` will remove duplicates of tuples corresponding to the `col list`
- You can only apply `DISTINCT` at the start of a list of columns
- So:
 - `SELECT A, DISTINCT B ...` is not permitted
 - But `SELECT COUNT(DISTINCT A) ...` is OK
 - Count of number of distinct values of A

SQL



- **So far: Basic Single-Table DML queries**
 - **SELECT (with DISTINCT)/FROM/WHERE**
 - **Aggregation: GROUP BY, HAVING**
 - **Presentation: ORDER BY, LIMIT**
- Extending basic SELECT/FROM/WHERE
 - Multi-table queries: JOINS
 - Aliasing in FROM and SELECT
 - Expressions in SELECT
 - Expressions, string comparisons, connectives in WHERE
 - Extended JOINS
 - The use of NULLS
- Query Composition
 - Set-oriented operations
 - Nested queries
 - Views
 - Common table expressions

Lots to cover!
Use vitamins and sections
to dig deeper.

SQL DML 1:

Basic Single-Table Queries



- **SELECT** [**DISTINCT**] *<column expression list>*
FROM *<single table>*
[**WHERE** *<predicate>*]
[**GROUP BY** *<column list>*
[**HAVING** *<predicate>*]]
[**ORDER BY** *<column list>*]
[**LIMIT** *<integer>*];

Conceptual Order of Evaluation



(5) **SELECT** [**DISTINCT**] *<col exp. list>*

(1) **FROM** *<single table>*

(2) [**WHERE** *<predicate>*]

(3) [**GROUP BY** *<column list>*]

(4) [**HAVING** *<predicate>*]]

(6) [**ORDER BY** *<column list>*]

(7) [**LIMIT** *<integer>*];

Will omit ORDER BY and LIMIT for now since they are primarily for presentation

SQL DML 1: Basic Single-Table Queries

Conceptual Order of Evaluation



- (5) **SELECT** [**DISTINCT**] <col exp. list> ➡ remove (project) cols not found in list, then remove dupl. rows
- (1) **FROM** <single table> ➡ for each tuple in table
- (2) [**WHERE** <predicate>] ➡ remove tuples that don't satisfy predicate (selection condition)
- (3) [**GROUP BY** <column list> ➡ form groups and perform all necessary aggregates per group
- (4) [**HAVING** <predicate>] ➡ remove groups that don't satisfy predicate

Q: Which aggregates are necessary?

A: All the aggregates that will be referred to in the **HAVING** or **SELECT** clause

Remember: this is all **conceptual** — actual approach for execution may be very different. But will provide the same result as this conceptual approach.

Putting it all together



- **SELECT** S.dept, **AVG**(S.gpa), **COUNT**(*)
FROM Students **AS** S
WHERE S.state = 'MA'
GROUP BY S.dept
HAVING **MAX**(S.gpa) >= 2
ORDER BY S.dept;
- Students (name, dept, gpa, state)
 - Start with all **tuples** in Students
 - **Throw** away those that aren't from MA
 - **Group by** S.dept, **compute aggregates** MAX(S.gpa), AVG(S.gpa), COUNT(*)
 - **Throw** away groups that don't have MAX(S.gpa)>=2
 - **Retain** only S.dept, AVG(S.GPA), COUNT(*)
 - **Order by** S.dept

Multi-Table Queries: Joins



- SELECT [DISTINCT] *<column expression list>*
FROM *<table1 [AS t1], ... , tableN [AS tn]>*
[WHERE *<predicate>*]
[GROUP BY *<column list>*[HAVING *<predicate>*]]
[ORDER BY *<column list>*];

SQL DML 1: Basic Single-Table Queries

Conceptual Order of Evaluation



Let's not worry about GROUP BY and HAVING for now, back to good old SELECT-FROM-WHERE
Extending it to GROUP BY and HAVING is straightforward (as is ORDER BY and LIMIT)

- (5) **SELECT** [**DISTINCT**] *<col exp. list>* ➡ *remove (project out) cols not found in list, then remove duplicate rows*
- (1) **FROM** *<table1><table2>...* ➡ *for each combinations of tuples in cross product of tables*
- (2) [**WHERE** *<predicate>*] ➡ *remove tuple combinations that don't satisfy predicate (selection condition)*
- (3) [**GROUP BY** *<column list>*] ➡ *form groups and perform all necessary aggregates per group*
- (4) [**HAVING** *<predicate>*] ➡ *remove groups that don't satisfy predicate*

Another way to think about a multi-table query is a query on a new relation that is the cross-product of tables in the FROM clause.

This is likely a really bad way to evaluate this query! We will discuss better ways subsequently.

Cross (Cartesian) Product



- **FROM** clause: all pairs of tuples, concatenated

Sailors

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...

Find sailors who've reserved
a boat



```
SELECT S.sid, S.sname, R.bid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...

Find sailors who've reserved
a boat cont



```
SELECT S.sid, S.sname, R.bid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	bid
1	Popeye	102
1	Popeye	101
2	OliveOyl	102

Table Aliases and Column Name Aliases



```
SELECT Sailors.sid, sname, bid
FROM Sailors, Reserves
WHERE Sailors.sid = Reserves.sid
```

Relation (range) variables (Sailors, Reserves) help refer to columns that are shared across relations.

We can also rename relations and use **new variables** ("AS" is optional for FROM)

```
SELECT S.sid, sname, bid
FROM Sailors AS S, Reserves AS R
WHERE S.sid = R.sid
```

We can also **rename attributes** too!

```
SELECT S.sid AS sailordid, sname AS sailordname, bid AS boatid
FROM Sailors AS S, Reserves AS R
WHERE S.sid = R.sid
```

More Aliases: Self-Joins



```
SELECT x.sname AS sname1,  
       x.age AS age1,  
       y.sname AS sname2,  
       y.age AS age2  
FROM Sailors AS x, Sailors AS y  
WHERE x.age > y.age
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sname1	age1	sname2	age2
Popeye	22	Bob	19
OliveOyl	39	Popeye	22
OliveOyl	39	Garfield	27
OliveOyl	39	Bob	19
Garfield	27	Popeye	22
Garfield	27	Bob	19

- Query for **pairs** of sailors where one is older than the other
- **Table** aliases in the **FROM** clause
 - Needed when the same table used multiple times (“self-join”)

Arithmetic Expressions



- SELECT S.age, **S.age-5** AS age1, **2*S.age** AS age2
FROM Sailors AS S
WHERE S.sname = 'Popeye'
- SELECT S1.sname AS name1, S2.sname AS name2
FROM Sailors AS S1, Sailors AS S2
WHERE **2*S1.rating = S2.rating - 1**

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

String Comparisons



- Old School SQL

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sname LIKE 'B_%'
```

`_` = any single char; `%` = zero or more chars

Returns Bob

- Standard Regular Expressions

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sname ~ 'B.*'
```

`.` = any char; `*` = repeat (zero or more instances of previous)

Note: can match anywhere in the string

Returns Bob and McBob

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19
5	McBob	3	35

SQLite note: `~` not supported.

Boolean Connectives



Sid's of sailors who reserved a red **OR** a green boat

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
        (B.color='red' OR B.color='green')
```

Boats

bid	bname	color
102	Titanic	green
101	Lusitania	red
100	Mayflower	orange

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
1	100	10/01

SQL



- So far: Basic Single-Table DML queries
 - SELECT (with DISTINCT)/FROM/WHERE
 - Aggregation: GROUP BY, HAVING
 - Presentation: ORDER BY, LIMIT
- Extending basic SELECT/FROM/WHERE
 - Multi-table queries: JOINS
 - Aliasing in FROM and SELECT
 - Expressions in SELECT
 - Expressions, string comparisons, connectives in WHERE
 - **Extended JOINS**
 - The use of NULLS
- Query Composition
 - Set-oriented operations
 - Nested queries
 - Views
 - Common table expressions

Join Variants



```
SELECT <column expression list>
FROM table_name
    [INNER | NATURAL
      | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
ON <qualification_list>
WHERE ...
```

- INNER is default
 - Same thing as what we've done so far, offers no additional convenience
 - Just present as a contrast to NATURAL and OUTER

Reminder



- Turn on video if you can
 - Turn off audio except when speaking
 - Don't do anything you wouldn't do normally
-
- Vitamin 1 deadline has been pushed
 - Project 1 should still be on track

Inner/Natural Joins



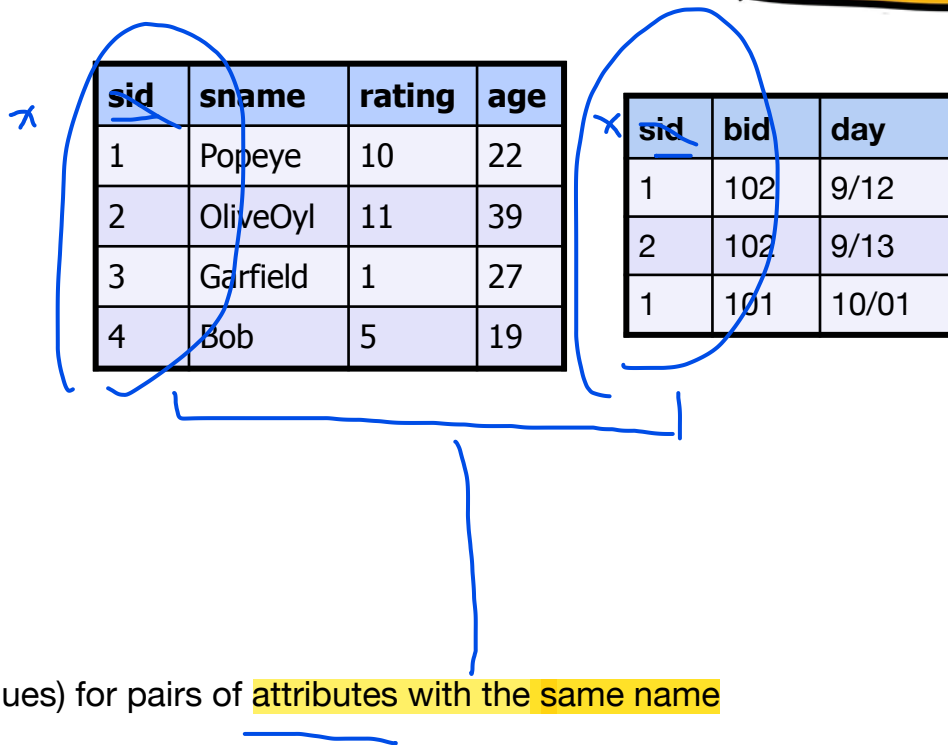
```
SELECT s.sid, s.sname, r.bid
FROM Sailors s, Reserves r
WHERE s.sid = r.sid
    AND s.age > 20;
```

```
SELECT s.sid, s.sname, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
WHERE s.age > 20;
```

```
SELECT s.sid, s.sname, r.bid
FROM Sailors s NATURAL JOIN Reserves r
WHERE s.age > 20;
```

- **ALL 3 ARE EQUIVALENT!**

- “NATURAL” means “**equi-join**” (i.e., identical values) for pairs of **attributes with the same name**



Left Outer Join



- Returns **all matched rows**, and **preserves all unmatched rows** from the table on the **left** of the join clause
 - (use **NULLs** in fields of **non-matching tuples**)
 - We'll talk about NULLs in a bit, but for now, think of it as N/A

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid;
```

Returns all sailors & bid for boat in any of their reservations

Note: no match for s.sid? r.bid IS NULL!

(3, Garfield, NULL) (4, Bob, NULL) in output

L

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

R

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

Right Outer Join



- Returns all matched rows, and preserves all unmatched rows from the table on the right of the join clause
 - (use NULLs in fields of non-matching tuples)

```
SELECT r.sid, b.bid, b.bname  
FROM Reserves r RIGHT OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats and sid for any sailor associated with the reservation.

Note: no match for b.bid? r.sid IS NULL!

Full Outer Join



- Returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.bname
FROM Reserves r FULL OUTER JOIN Boats b
ON r.bid = b.bid
```

- Returns all boats & all information on reservations
- No match for r.bid?
 - b.bid IS NULL AND b.bname IS NULL!
- No match for b.bid?
 - r.sid IS NULL!

SQLite note: RIGHT/FULL OUTER JOIN not supported.

Brief Detour: NULL Values



- Values for any data type can be **NULL**
 - Indicates the value is present but unknown or is inapplicable
 - Also comes naturally from Outer joins
- The presence of **null complicates** many issues. E.g.:
 - **Selection predicates** (WHERE)
 - **Aggregation**

NULL in the WHERE clause



```
SELECT * FROM sailors  
WHERE rating > 8;
```

Q: Should Popeye be in the output?

Not really.

Likewise for

```
SELECT * FROM sailors  
WHERE rating <= 8;
```

sid	sname	rating	age
1	Popeye	NULL	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

NULL in the WHERE clause



```
SELECT * FROM sailors
WHERE rating > 8 OR rating <= 8;
```

sid	sname	rating	age
1	Popeye	NULL	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

This is really funky — we have a tautology in the WHERE clause, but Popeye will still not be output

To force certain outputs can use IS NULL or IS NOT NULL conditions

```
SELECT * FROM sailors
WHERE rating > 8 OR rating <= 8 OR rating IS NULL;
```

This will correctly output all tuples in this setting

More generally, we need an extension to Boolean logic to support this

Correctly Reasoning about NULLs



- Several Ingredients:
 - We need a way to evaluate **unit predicates**, a way to combine them, and a way to decide whether to output
- Ingredient 1: Evaluating unit predicates
 - **(x op NULL) evaluates to NULL (IDK!)**

```
SELECT 100 = NULL;  
SELECT 100 < NULL;
```
 - **IS NULL** evaluates to **True** if NULL, **False** otherwise
- Ingredient 3: Deciding to output
 - When the **WHERE** evaluates to **NULL**, **do not output the tuple**

```
SELECT * FROM sailors;  
SELECT * FROM sailors WHERE rating > 8;  
SELECT * FROM sailors WHERE rating <= 8;
```
- Ingredient 2: Combining predicates
 - **Three-valued logic**, an extension of **two-valued (Boolean) logic**

NULL in Boolean Logic

Three-valued logic: truth tables!

AND	T	F	N
T	T	F	N
F	F	F	F
N	N	F	N

OR	T	F	N
T	T	T	T
F	T	F	N
N	T	N	N

Let's build intuition by going through examples

`SELECT * FROM sailors WHERE rating > 8 OR rating <= 8; \Rightarrow N`

`SELECT * FROM sailors WHERE NOT (rating > 8); $\Rightarrow \neg N \Rightarrow N$`

`SELECT * FROM sailors WHERE rating > 8 OR TRUE;`

NOT	T	F	N
	F	T	N

General rule: **NULL values** are treated as **"I Don't Know"** – can be either **true or false**

NULL and Aggregation



General rule: NULL ****column values**** are **ignored** by **aggregate functions**

```
SELECT count(*) FROM sailors;
```

```
SELECT count(rating) FROM sailors;
```

```
SELECT sum(rating) FROM sailors;
```

```
SELECT avg(rating) FROM sailors;
```

NULL and Aggregation



General rule: NULL **column values are ignored by aggregate functions**

```
SELECT count(*) FROM sailors; // count sailors
```

```
SELECT count(rating) FROM sailors; // count sailors with non-NULL ratings
```

```
SELECT sum(rating) FROM sailors; // sum of non-NULL ratings
```

```
SELECT avg(rating) FROM sailors; // avg of non-NULL ratings
```

NULLs: Summary



- **NULL op x**; x op NULL is NULL
- **WHERE NULL**: do not send to output
- Boolean connectives: **3-valued logic**
- **Aggregates** ignore **NULL-valued** inputs

SQL



- Basic Single-Table DML queries
 - SELECT (with DISTINCT)/FROM/WHERE
 - Aggregation: GROUP BY, HAVING
 - Presentation: ORDER BY, LIMIT
- Extending basic SELECT/FROM/WHERE
 - Multi-table queries: JOINS
 - Aliasing in FROM and SELECT
 - Expressions in SELECT
 - Expressions, string comparisons, connectives in WHERE
 - Extended JOINS
 - The use of NULLS
- **Query Composition**
 - **Set-oriented operations**
 - **Nested queries**
 - **Views**
 - **Common table expressions**

Let's talk about Sets and Bags



- **Set** = no duplicates {🍏, 🍏, 🍊, 🍋, 🍌}
- **Bag / Multi-set** = duplicates allowed {🍏, 🍏, 🍏, 🍊, 🍊, 🍊}
- As we saw earlier SQL uses bag semantics
 - That is, there can be **multiple** copies of each tuple in a relation
- How do we “add/subtract” tuples across relations?
 - We can do so **operators** that enforce either bag or set-based semantics

Operators with Set Semantics



- Set: a collection of distinct elements
 - In the relational parlance: each tuple/row is unique
- Ways of manipulating/combining sets
 - A UNION B: distinct tuples in A or B
 - A INTERSECT B: distinct tuples in A and B
 - A EXCEPT B: distinct tuples in A but not in B
- Basically, we treat tuples within a relation as elements of a set

Using Set Semantics with SQL



Note: R and S are relations. They are not sets, since they have duplicates.

Assume these are all tuples: A, B, C, D, E

$R = \{A, A, A, A, B, B, C, D\}$
 $S = \{A, A, B, B, B, C, E\}$

- **UNION**

$\{A, B, C, D, E\}$

- **INTERSECT**

$\{A, B, C\}$

- **EXCEPT**

$\{D\}$

Reserves

sid	bid	day
1	102	9/12
1	102	9/12
2	101	10/01

Q: What does
(SELECT * FROM Reserves)
UNION
(SELECT * FROM Reserves)
give us?

“ALL”: Multiset Semantics



$R = \{\text{A, A, A, A, B, B, C, D}\} = \{\text{A(4), B(2), C(1), D(1)}\}$

$S = \{\text{A, A, B, B, B, C, E}\} = \{\text{A(2), B(3), C(1), E(1)}\}$

“UNION ALL”: Multiset Semantics



$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

• UNION ALL: **sum of cardinalities**

$\{A(4+2), B(2+3), C(1+1), D(1+0), E(0+1)\}$
 $= \{A, A, A, A, A, A, B, B, B, B, B, C, C, D, E\}$

Reserves

sid	bid	day
1	102	9/12
1	102	9/12
2	101	10/01

4 \Rightarrow

2 \Rightarrow

Q: What does
(SELECT * FROM Reserves)
UNION ALL
(SELECT * FROM Reserves)
give us?

2 1
2 1
—
4 2

“INTERSECT ALL”: Multiset Semantics



$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

• INTERSECT ALL: min of cardinalities

$\{A(\min(4,2)), B(\min(2,3)), C(\min(1,1)),$
 $D(\min(1,0)), E(\min(0,1))\}$
 $= \{A, A, B, B, C\}$

“EXCEPT ALL”: Multiset Semantics



$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- EXCEPT ALL: difference of cardinalities
 $\{A(4-2), B(2-3), C(1-1), D(1-0), E(0-1)\}$
 $= \{A, A, D\}$

Set/Bag Operators



- A **UNION** B, A **INTERSECT** B, A **EXCEPT** B perform set-based operations treating tuples in A and B as **sets**
- A **UNION ALL** B, A **INTERSECT ALL** B, A **EXCEPT ALL** B perform **bag-based** operations treating tuples in A and B as bags
- **Note:** for these operations to be applied correctly, the schema for A and B must be the same!

Combining Predicates



- Subtle connections between:
 - Boolean logic in WHERE (i.e., AND, OR)
 - Set operations (i.e. INTERSECT, UNION)
- Let's see some examples...

Sid's of sailors who reserved a red **OR** a green boat



```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'
```

UNION

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='green'
```

VS...

```
SELECT DISTINCT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND  
      (B.color='red' OR B.color='green')
```

These two give the exact same result!

HW:

- a) What if we did **UNION ALL** instead?
- b) What if we omitted **DISTINCT**?

Sid's of sailors who reserved a red **AND** a green boat



```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'
```

INTERSECT

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='green'
```

VS...

```
SELECT DISTINCT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND (B.color='red' AND B.color='green')
```

The first query works fine... but the second query doesn't work. Why?

SQL



- Basic Single-Table DML queries
 - SELECT (with DISTINCT)/FROM/WHERE
 - Aggregation: GROUP BY, HAVING
 - Presentation: ORDER BY, LIMIT
- Extending basic SELECT/FROM/WHERE
 - Multi-table queries: JOINS
 - Aliasing in FROM and SELECT
 - Expressions in SELECT
 - Expressions, string comparisons, connectives in WHERE
 - Extended JOINS
 - The use of NULLS
- Query Composition
 - Set-oriented operations
 - **Nested queries**
 - Views
 - Common table expressions

Query Composition



- We've already seen one way of combining results across multiple queries via set and bag-based operations
- Now, we'll talk about “nesting” queries inside other queries
 - Nesting and subqueries
 - Views to refer to frequent query expressions
 - Common Table Expressions

Nested Queries: IN



- *Names of sailors who've reserved boat #102:*

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN
      (SELECT R.sid
       FROM   Reserves R
       WHERE  R.bid=102)
```

subquery



Here, the results of this subquery are treated as a (multi)set, with membership of S.sid checked in the set using the IN operator

Nested Queries: NOT IN



- *Names of sailors who've **not** reserved boat #103:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid NOT IN
        (SELECT  R.sid
         FROM    Reserves R
         WHERE   R.bid=103)
```


Nested Queries with Correlation



- So far, we've studied ways to nest query results and treat it as a “set” with membership in the set checked
 - using ... ***val [NOT] IN (nested query)***
- We can also check if a nested query result is empty/not
 - using ... ***[NOT] EXISTS (nested query)***
- *Names of sailors who've reserved boat #102:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS
        (SELECT *
         FROM Reserves R
         WHERE R.bid=102 AND S.sid=R.sid)
```

- Correlated subquery is **conceptually** recomputed for each Sailors tuple.

More on Set-Comparison Operators



- We've seen: [NOT] IN, [NOT] EXISTS
- Other forms: op ANY, op ALL

Find sailors whose rating is greater than that of *some* sailor called Popeye:

```
SELECT *
FROM   Sailors S
WHERE  S.rating > ANY
      (SELECT S2.rating
       FROM   Sailors S2
       WHERE  S2.sname='Popeye')
```

SQLite note: ANY/ALL not supported.

A Tough One: “Division”



- Relational Division: “Find sailors who’ve reserved all boats.”
Said differently: “Sailors with no missing boats”

```
SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS
```

For S, this is the set of all boats
they have not reserved

```
(SELECT B.bid  
FROM Boats B  
WHERE NOT EXISTS (SELECT R.bid  
FROM Reserves R  
WHERE R.bid=B.bid  
AND R.sid=S.sid ))
```

For S and B, this is the set
of reservations of B for S

ARGMAX?



- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating >= ALL  
      (SELECT S2.rating  
       FROM   Sailors S2)
```

VS

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating =  
      (SELECT MAX(S2.rating)  
       FROM   Sailors S2)
```

These are exactly the same!

ARGMAX?



- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating >= ALL  
      (SELECT S2.rating  
       FROM   Sailors S2)
```

VS

```
SELECT *  
FROM   Sailors S  
ORDER BY rating DESC  
LIMIT 1;
```

These are not the same if there are multiple such Sailors

Views: Named Queries



```
CREATE VIEW view_name AS select_statement
```

- Makes development simpler, convenient
- Often used for security
- Not “materialized” [but there are materialized views as well!]

```
// Counts of reservations for red colored boats
```

```
CREATE VIEW Redcount AS
```

```
    SELECT B.bid, COUNT(*) AS scount  
    FROM Boats B, Reserves R  
    WHERE R.bid=B.bid AND B.color='red'  
    GROUP BY B.bid
```

Views Instead of Relations in Queries



```
CREATE VIEW Redcount AS
SELECT B.bid, COUNT(*) AS scount
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid;
```

bid	scount
102	1

```
SELECT * from Redcount WHERE scount<10;
```

Subqueries in FROM



Like a “view on the fly”!

```
SELECT *  
FROM  
  (SELECT B.bid, COUNT (*)  
   FROM Boats B, Reserves R  
   WHERE R.bid = B.bid AND B.color = 'red'  
   GROUP BY B.bid) AS Redcount(bid, scout)  
WHERE scout < 10
```


WITH a.k.a. common table expression (CTE)



Another “view on the fly” syntax:

```
WITH Redcount(bid, scout) AS  
(SELECT B.bid, COUNT (*)  
FROM Boats B, Reserves R  
WHERE R.bid = B.bid AND B.color = 'red'  
GROUP BY B.bid)
```

```
SELECT * FROM Reds  
WHERE scout < 10
```

Can have many queries in WITH



Cascade of queries: Redcount -> UnpopularReds

```
WITH Redcount(bid, scount) AS  
(SELECT B.bid, COUNT (*)  
FROM Boats B, Reserves R  
WHERE R.bid = B.bid AND B.color = 'red'  
GROUP BY B.bid),
```

```
UnpopularReds AS  
(SELECT *  
FROM Redcount  
WHERE scount < 10)
```

```
SELECT * FROM UnpopularReds;
```

ARGMAX GROUP BY?



- More complex variation of previous argmax
- Find the sailors with the highest rating per age

```
WITH maxratings(age, maxrating) AS  
  (SELECT age, max(rating)  
   FROM Sailors  
   GROUP BY age)
```

```
SELECT S.*  
  FROM Sailors S, maxratings m  
 WHERE S.age = m.age  
       AND S.rating = m.maxrating;
```

Testing SQL Queries



- Typically not every database instance will reveal every bug in your query.
 - Eg: database instance without any rows in it!
- Best to try to reason about behavior across all instances
- Also helpful: constructing test data.

Tips for Generating Test Data



- Generate **random data**
 - e.g. using a service like mockaroo.com
- Try to construct data that could check for the following potential errors:
 - Incorrect output schema
 - Output may be missing rows from the correct answer (false negatives)
 - Output may contain incorrect rows (false positives)
 - Output may have the wrong number of duplicates.
 - Output may not be ordered properly.

Summary



- You've now seen SQL—you are armed.
- A declarative language
 - Somebody has to translate to algorithms though...
 - The RDBMS implementor ... i.e. you!

Summary Cont



- The data structures and algorithms that make SQL possible also power:
 - NoSQL, data mining, scalable ML, network routing...
 - A toolbox for scalable computing!
 - Start talking about that in the next set of slides!
- We skirted questions of good database (schema) design
 - a topic we'll consider in greater depth later