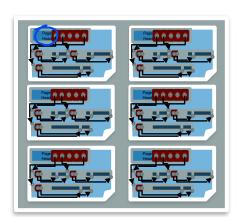
# Disk Representations: Files, Pages, Records



#### Overview: Files of Pages of Records



- Overall:
  - Each table is stored in one or more OS files
  - Each file contains many pages
  - Each page contains contains many records + wple
- Pages are the common currency understood by multiple layers:
  - Managed on disk by the disk space manager: pages read/written to physical disk/files
  - Managed in memory by the buffer manager: higher levels of DBMS only operate in memory



### Files of Pages of Records



- Let's talk about a single table for now
- DB FILE: A collection of pages, each containing a collection of records.
- API for higher layers of the DBMS:
  - Reads:
    - Fetch a particular record by record id ...
      - Record id is a pointer encoding pair of (pageID, location on page)
    - Scan all records
      - Possibly with some conditions on the records to be retrieved
  - Updates: Insert/delete/modify record
- This abstraction could span multiple OS files and even machines

### Many DB File Structures



Information is stored in files in multiple different ways

- Unordered Heap Files
  - Records placed arbitrarily across pages
- Clustered Heap Files
  - Records and pages are grouped in some meaningful way
- Sorted Files
  - Pages and records are in strict sorted order
- Index Files
  - B+ Trees, Linear Hashing, ...
  - May contain records or point to records in other files
- Focus on Unordered Heap Files for now...

#### Unordered Heap Files



- Collection of records in no particular order
  - Not to be confused with "heap" data-structure: efficient max/min
- As file shrinks/grows, pages (de)allocated
- To support record level operations, we must
  - Keep track of the pages in a file
  - Keep track of free space on pages
  - Keep track of the records on a page

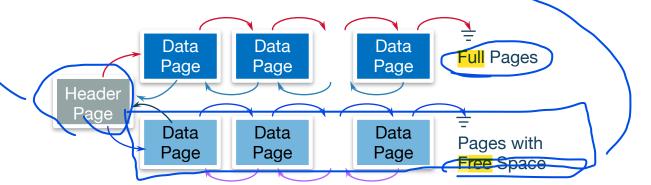
### Take 1: Heap File as List



- Heap file has one special "Header page"

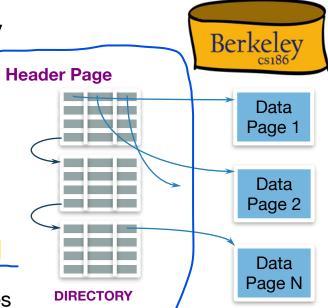
  Schema. Ref

  Location of the heap file and the header page saved e.g., in catalog
- Each page contains 2 "pointers" plus free space and data
- What is wrong with this?
  - How do I find a page with enough space for a 20 byte record
  - A: Need to access many pages (w/ free space) to check



Take 2: Use a Page Directory

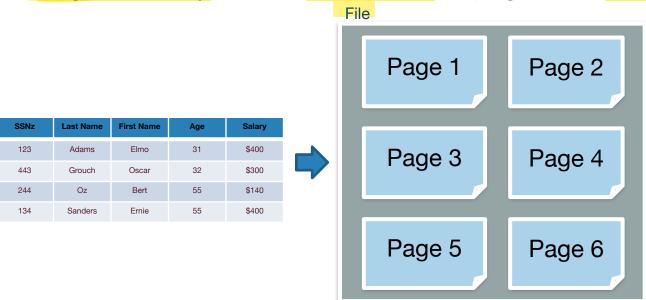
- Directory, with multiple Header Pages, each encoding:
  - A pointer to page
  - #free bytes on the page
- There can be multiple such header pages
- Header pages accessed often → likely in cache
- Finding a page to fit a record required far fewer page loads than linked list. Why?
  - One header page load reveals free space of many pages
- You can imagine optimizing the page directory further
  - E.g., compressing header page, keeping header page in sorted order based on free space, etc.
  - But diminishing returns?



# Summary



- Table encoded as files which are collections of pages
- Page directory provides locations of pages and free space





#### **PAGE LAYOUT**

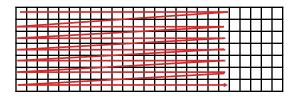
#### A Note On Imagery



Data (in memory or disk) is stored in linear order



- This doesn't fit nicely on screen
  - So we will "wrap around" the linear order into a rectangle



#### Page Basics: The Header



- Header may contain "metadata" about the page, e.g.
  - Number of records
  - Free space
  - Maybe a next/last pointer
  - Bitmaps, Slot Table
  - (We'll talk about why all of these later)

Page Header

### Things to Address



#### Some options:

- Strings
- Record length? Fixed or Variable
- Page layout? Packed or Unpacked

free space

#### Some questions:

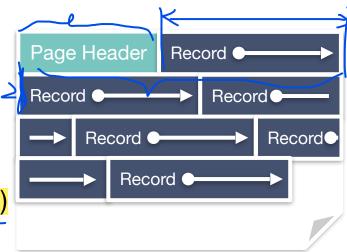
- Find records by record id?
  - Record id (Page, Location in Page)
- How do we add and delete records?

Page Header

#### Fixed Length Records, Packed



- Pack records densely
- Record id = (pageId, "location in page")?
  - (pageld, record number in page)!
  - We know the offset from start of page!
    - Offset = header + (record size) x (n-1)
- Easy to add: just append
- Delete?

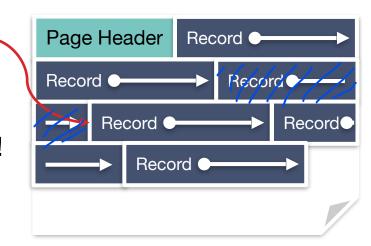


### Fixed Length Records, Packed, Pt 2.

Record id:



- Pack records densely
- Record id = (pageId, "location in page")?
  - (pageld, record number in page)!
  - We know the offset from start of page!
- Easy to add: just append
- Delete?
  - Say we delete (Page 2, Record 3)

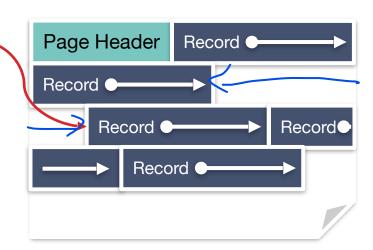


#### Fixed Length Records: Packed, Pt 3.

Record id:



- Pack records densely
- Record id = (pageId, "location in page")?
  - (pageId, record number in page)!
  - We know the offset from start of page!
- Easy to add: just append
- Delete?
  - Say we delete (Page 2, Record 3)
  - Now free space... need to reorg

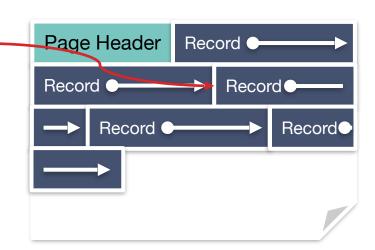


# Fixed Length Records: Packed, Pt. 5

Record id:



- Pack records densely
- Record id = (pageId, "location in page")?
  - (pageld, record number in page)!
  - We know the offset from start of page!
- Easy to add: just append
- Delete?
  - Packed implies re-arrange!
  - "record id" (Page 2, Record 4) now need to be updated to (Page 2, Record 3)
  - Record Ids need to be updated!
    - Could be expensive if they're in other files.

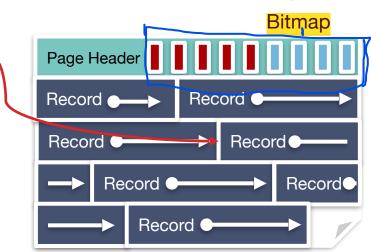


# Fixed Length Records: Unpacked



Record id: (Page 2, Record 4)

- Bitmap denotes "slots" with records
- Record id = (pageId, "location in page")?
  - (pageld, slotld)
- Insert: find first empty slot in bitmap
- Delete: ?



# Fixed Length Records: Unpacked, Pt. 2 Berkeley

Record id:



Bitmap denotes "slots" with records

Record id = (pageld, "location in page")?

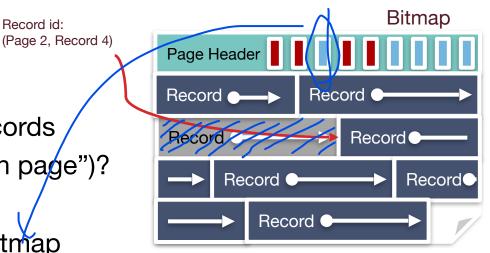
(pageld, slotld)

**Insert**: find first empty slot in bitmap

**Delete:** clear bit

No reorganization needed!

Small cost of a bitmap, which can be very compact

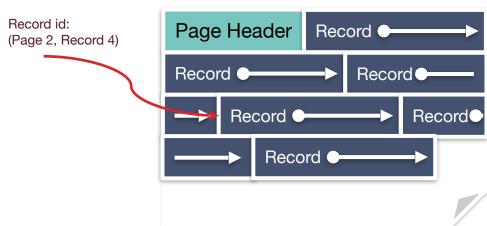


#### Variable Length Records



 We've already seen that packed isn't the best idea, so let's consider the unpacked case

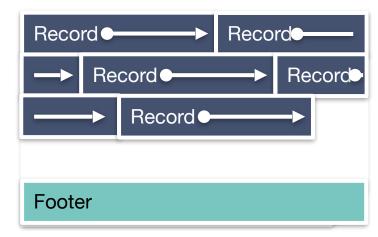
- How do we know where each record begins (mapping recordid to location)?
- What happens when we add and delete records?



#### First: Relocate metadata to footer

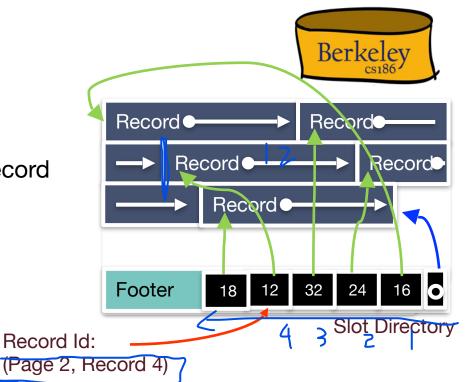


We'll see why this is handy shortly...



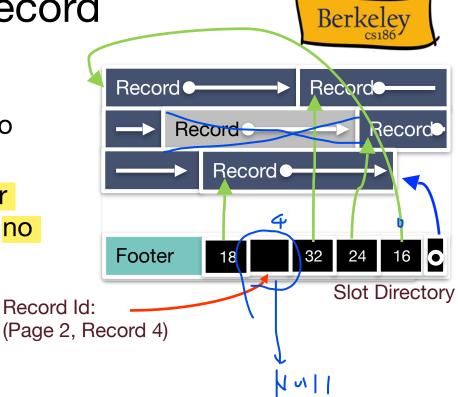
#### Slotted Page

- Introduce slot directory in footer
  - Pointer to free space
  - Length + Pointer to beginning of record
    - reverse order
- Record ID = location in slot table
  - from right
- Delete?
  - e.g., 4th record on the page



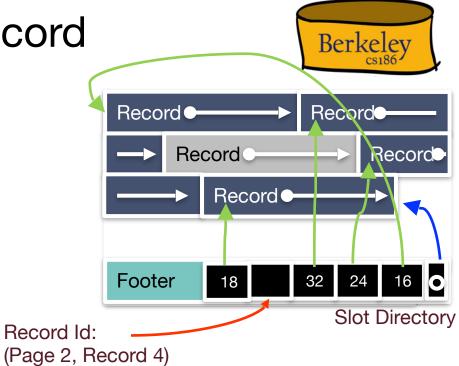
Slotted Page: Delete Record

- Delete record (Page 2, Record 4):
  - Set 4th slot directory pointer to null
  - Doesn't affect pointers to other records (no internal reorg, and no updating of external pointers)



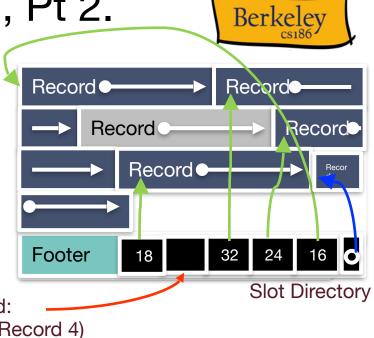
Slotted Page: Insert Record

Insert:



Slotted Page: Insert Record, Pt 2.

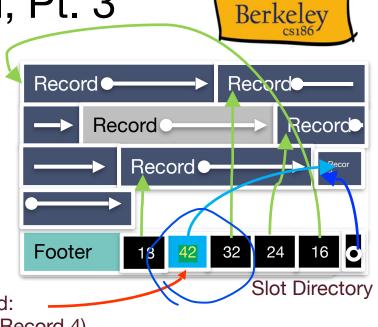
- Insert:
  - Place record in free space on page



Record Id:

Slotted Page: Insert Record, Pt. 3

- Insert:
  - Place record in free space on page
  - Create pointer length pair in next open slot in slot directory

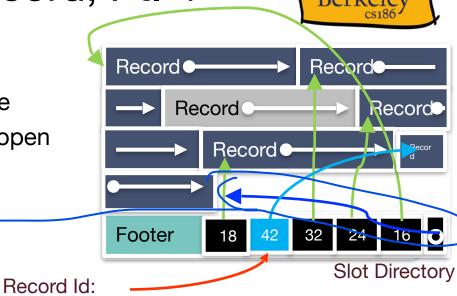


Record Id:

Slotted Page: Insert Record, Pt. 4

Berkeley cs186

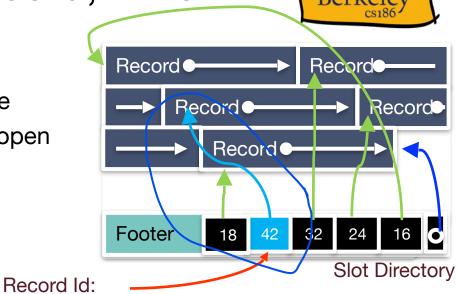
- Insert:
  - Place record in free space on page
  - Create pointer/length pair in next open slot in slot directory
  - Update the free space pointer
  - Fragmentation?



Slotted Page: Insert Record, Pt. 6

Berkeley cs186

- Insert:
  - Place record in free space on page
  - Create pointer/length pair in next open slot in slot directory
  - Update the free space pointer
  - Fragmentation?
    - Reorganize data on page!



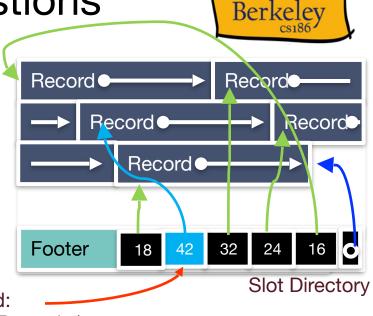
Slotted Page: Leading Questions

Berkelev

- Reorganize data on page
  - Is this safe?
    - Yes this is safe because records ids don't change. Record ids refer to slots
- When should I reorganize?
  - We could re-organize on delete
  - Or wait until fragmentation blocks record addition and then reorganize! or 2 y
  - Often pays to be a little sloppy if page never gets more records.

Record Id: (Page 2, Record 4)

- What if we need more slots?
  - Let's see...

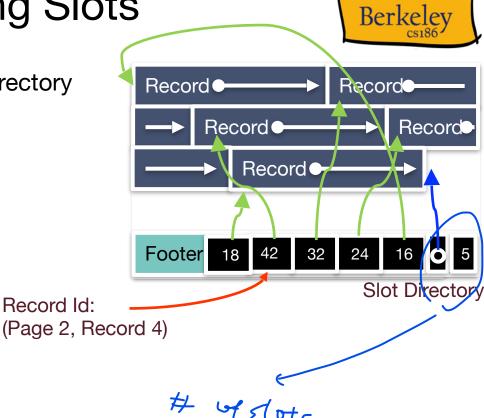


Slotted Page: Growing Slots

Record Id:

Tracking **number of slots** in slot directory

Empty or full

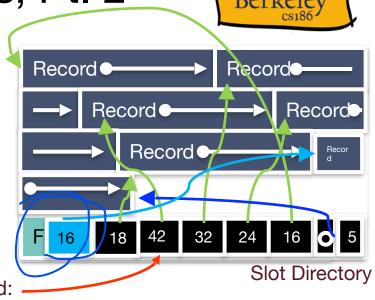


Slotted Page: Growing Slots, Pt. 2

Berkeley CS186

- Tracking number of slots in slot directory
  - Empty or full
- If full slots = number of slots, then extend slot directory
- To extend slot directory
  - Slots grow from end of page inward

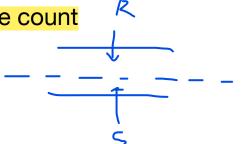
  - Easy!

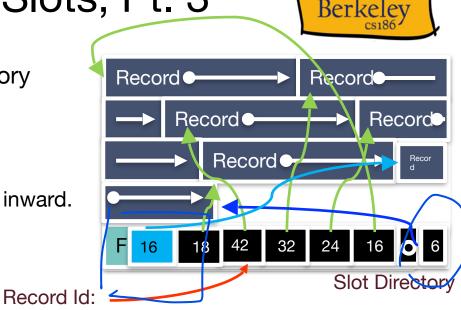


Slotted Page: Growing Slots, Pt. 3

Berkelev

- Tracking number of slots in slot directory
  - Empty or full
- Extend slot directory
  - Slots grow from end of page inward
  - Records grow from beginning of page inward.
  - Easy!
- And update count





### Slotted Page: Summary

- Typically use Slotted Page
  - Good for variable and fixed length records
- Not bad for fixed length records too.
  - Why?
  - Fixed length records also have NULL fields

Slot Directory

Record

Recor

Record

NULL values can be "squashed" and indicated using a flag, avoiding full attribute length storage

Record •

Record •

Record •

32

24

 But, if we have only non-NULL fields, can be worth the optimization of fixed-length format



#### **RECORD LAYOUT**

#### Record Formats



- Each record in a table/relation has a fixed combo of types
- Relational databases also use same page format for data on disk or in memory

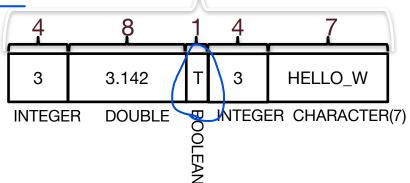
schema

- Save cost of conversion (known as serialization/deserialization)
- Assume System Catalog stores the Schemat
  - No need to store type information with records (save space!)
  - Catalog is just another table
- Goals:
  - Fast access to fields (why?)
  - Records should be compact
- Easy Case: Fixed Length Fields
- Interesting Case: Variable Length Fields

#### Record Formats: Fixed Length



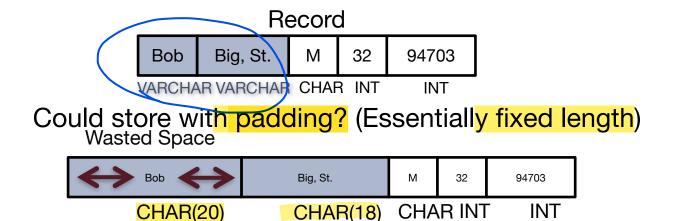
- Finding i'th field?
  - done via arithmetic (fast)
- Making it more compact?
  - If all fields are not-null, no good way of compacting
  - Else apply variable length techniques, next



#### Record Formats: Variable Length

What happens if fields are variable length?



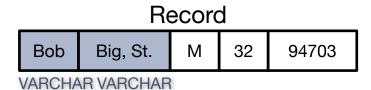


- But have to account fo<mark>r largest possible string (wasteful) or rearrange as soon as a larger string comes (inefficient).</mark>
- Could store with delimiters (e.g., commas)?
  But makes it hard to find fields and also ensure that commas are not
- part of the string

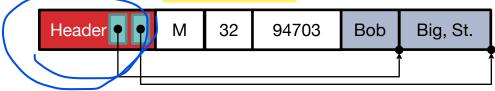
# Record Formats: Variable Length, Pt. 7



What happens if fields are variable length?



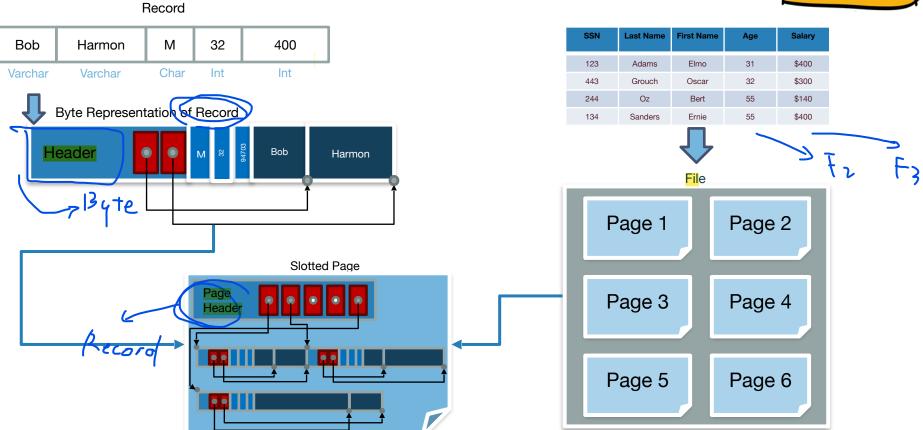
Solution: introduce a record header



- Easy access to fields, and almost as compact as can be (modulo header)
  - Same approach can be used to squash fixed length null fields w. many nulls

#### Overview: Representations





#### Files: Summary



- DBMS "File" contains pages, and records within pages
  - Heap files: unordered records organized with directories
- Page layouts
  - Fixed-length packed and unpacked
  - Variable length records in slotted pages, with intra-page reorg
- Variable length record format
  - Direct access to i'th field and null values